

Surfacing Suicidal Vulnerability Through Simulated Social Interaction: Per-Person Language Model Agents as Communicative Stress Tests

Abstract

Suicidal vulnerability may be encoded in everyday communication patterns but diluted in routine digital interactions. We introduce a method for surfacing this latent signal: training per-person language model agents on individuals’ real text messages and placing them in simulated social interactions—a “communicative stress test.” Using data from 79 adults with recent suicidal ideation, we fine-tuned individual low-rank adaptation (LoRA) adapters on an open-source language model (Qwen3-8B), then evaluated three paradigms: an all-pairs arena (2,485 paired conversations across 71 agents), a caption-trained comparison (agents trained on VLM-generated screenshot descriptions rather than messages), and a standardized probe arena (5 theory-driven personas $\times$ 71 agents). Message-trained agents whose real-world counterparts had higher ecological momentary assessment (EMA)-measured suicidal ideation produced significantly more suicide-relevant language (Spearman $\rho = .438$, $p < .001$; Cohen’s $d = 0.77$), with absolutist language as the strongest single predictor ($r = .607$, $p < .001$). The standardized probe arena was associated with suicidal ideation at $r = .576$ ($p < .001$), compared to $r = .477$ for the all-pairs arena and $r = .461$ for raw messages. Strikingly, a single neutral small-talk probe—with no emotional provocation—performed

nearly as well ($r = .551$), and combining just two probes (Neutral Acquaintance and Supportive Therapist) yielded the best cross-validated prediction (LOOCV $r = .570$), with additional probes adding noise rather than signal. Theory-driven construct-specific probe targeting largely failed: probes designed to elicit specific interpersonal constructs did not outperform the unstructured arena for their target outcomes. Caption-trained agents produced no clinical signal despite $50\times$ more training data, confirming that the predictive validity is specifically attributable to interpersonal communication patterns. Validation analyses confirmed that the clinical signal is person-specific: shuffling adapters across participants eliminated prediction ($r = .071$, $p = .577$), while the mismatched adapter’s output tracked the adapter-owner’s SI ($r = .426$, $p < .001$). A prompt ablation demonstrated that the signal partially survives without disclosure-encouraging language ($r = .430$ vs. $r = .576$), indicating that the prompt amplifies an existing signal rather than creating it. A no-LoRA floor control established that the base model produces near-zero between-person variance in risk language; the adapter’s contribution is individual differentiation, not mean-level risk. The method cannot model participants who produce minimal digital text (excluded message count $M = 12$ vs. included $M = 1,202$), a fundamental limitation of text-based approaches. Temporal split validation demonstrated moderate stability (Spearman-Brown = .538), with both halves independently predicting SI. These findings demonstrate that simulated social interaction can amplify latent vulnerability signals in a proof-of-concept that bridges digital phenotyping, generative AI, and suicide theory.

Keywords: suicidal ideation, large language models, generative agents,

1. Introduction

Suicide remains a leading cause of death globally, and predicting who is at risk continues to be one of the most difficult problems in clinical science (Franklin et al., 2017). Despite decades of research, clinicians perform only marginally better than chance at predicting suicidal behavior (Large et al., 2017), and traditional risk factor approaches have reached apparent ceiling effects in predictive accuracy (Belsher et al., 2019). A central challenge is that suicidal ideation fluctuates rapidly—on the order of hours (Kleiman et al., 2017)—and the interpersonal processes theorized to drive it (perceived burdensomeness, thwarted belonging; Joiner 2005; Van Orden et al. 2010) are inherently relational, manifesting most clearly in social interaction rather than in static self-report.

Two converging developments create new opportunities for studying these dynamics. First, digital phenotyping—the continuous measurement of behavior through smartphones and wearable sensors—has produced rich longitudinal records of how individuals communicate in daily life. Screenomics, which captures smartphone screenshots at five-second intervals (Reeves et al., 2021), provides near-complete records of text messages sent and received. When combined with ecological momentary assessment (EMA) of suicidal ideation, these data reveal moment-to-moment relationships between digital communication and suicide risk (Ammerman et al., 2026a). In adolescent populations, intensive longitudinal smartphone sensing has shown particular

promise: Auerbach et al. (2023) demonstrated that mood disturbances measured via smartphone predicted clinically significant suicidal ideation on a weekly basis, and subsequent work used GPS data to detect next-week suicide risk (Auerbach et al., 2025). However, analyzing raw message content or sensor data for suicide risk faces fundamental limitations: most messages are logistical and routine, risk-relevant content is sparse and diluted, and direct surveillance of private messages raises profound ethical concerns.

Second, the emergence of large language models (LLMs) capable of producing human-like text has enabled a new class of research: generative agent simulation. Park et al. (2023) demonstrated that LLM-powered agents placed in a shared environment exhibit emergent social behaviors without explicit programming. Subsequent work has scaled this to 1,000 agents parameterized from real interview data, achieving 85% behavioral fidelity (Park et al., 2024). More broadly, LLMs have been used as proxies for human participants—Argyle et al. (2023) introduced “silicon sampling” to simulate survey responses, Aher et al. (2023) replicated classic psychology experiments, and Peters & Matz (2024) showed that LLMs can infer Big Five traits from social media posts ($r = 0.29$; see also Brickman et al. 2025). Applications to mental health have begun to appear, including agent-based simulations of student wellbeing (Zhao et al., 2025), frameworks for modeling socio-environmental determinants (Kambeitz & Meyer-Lindenberg, 2025), and simulations of vulnerable users interacting with AI chatbots (Qiu et al., 2025). Louie et al. (2024) validated the use of expert-designed probe personas for patient simulation. Concurrently, there has been a rapid expansion of LLMs applied directly to suicide prevention (Holmes et al., 2025). However, Larooij &

Törnberg (2025b) offered a critical review arguing that LLMs may exacerbate rather than solve the validation challenges inherent in agent-based social
50 simulation.

A substantial psycholinguistic literature has established that language carries reliable markers of suicidal risk. Seminal work by Stirman & Pennebaker (2001) demonstrated that suicidal poets used more first-person singular pronouns and fewer first-person plural pronouns than non-suicidal poets. Al-Mosaiwi & Johnstone (2018) showed that absolutist thinking—operationalized
55 through words like “nothing,” “completely,” and “always”—is a cognitive-linguistic marker specific to suicidal ideation beyond what negative emotion words alone capture. More broadly, a systematic review (Ophir et al., 2022) identified consistent markers across studies, including elevated self-referential language, death-related vocabulary, and reduced future-oriented
60 words. Dictionary-based approaches have been central to translating these findings into applied risk detection, with the Linguistic Inquiry and Word Count (LIWC; Pennebaker et al. 2015) program applied to suicide notes (Pestian et al., 2012), crisis hotline transcripts (Huang et al., 2024), and
65 clinical interviews (Tausczik & Pennebaker, 2010). In clinical messaging contexts, Swaminathan et al. (2023) demonstrated that a two-stage NLP system combining keyword filtering with machine learning reduced crisis message triage time from 9 hours to 8–13 minutes. In our own prior work, we applied dictionary-based classification to text messages extracted from smartphone
70 screenshots, finding that within-person increases in negative emotional expressions co-occurred with elevated suicidal ideation (Ammerman et al., 2026a).

The present study bridges digital phenotyping and generative agent simulation by training per-person language model agents on individuals' actual
75 text messages and placing those agents in simulated social interactions to surface latent communication patterns associated with suicidal vulnerability (Figure 1). For each participant, we fine-tune a low-rank adaptation (LoRA) adapter on an open-source language model (Qwen3-8B) using their real messages, creating an agent that encodes their communicative style. We then
80 place each agent in standardized conversations with probe personas designed to elicit specific interpersonal dynamics, and score the generated language for risk-relevant content using the dictionary-based approach described above. We conceptualize this as a *communicative stress test*—analogous to cardiac stress testing, where a resting ECG may appear normal but underlying vul-
85 nerability becomes visible under exertion. In our paradigm, routine text messages are the resting state, and the simulated conversation is the treadmill. This approach shares a conceptual affinity with recent work using multi-agent AI systems to systematically probe wearable sensor data for clinically relevant biomarkers (Kim et al., 2025), though our method operates on language
90 rather than physiological signals and creates per-person models rather than population-level features.

This approach is novel relative to the existing agent simulation literature in a fundamental way. Prior methods rely on *prompt-based conditioning*—providing demographic descriptions (Argyle et al., 2023), interview tran-
95 scripts (Park et al., 2024), or expert-designed persona specifications (Louie et al., 2024)—to shape model behavior. This faces an individuation problem: Peng et al. (2025) found that digital twins created through prompting

correlate with real human responses at only $r = 0.20$, and Li et al. (2025) demonstrated that even state-of-the-art models struggle to maintain behavioral consistency. Our approach differs by *fine-tuning* per-person adapters on
100 actual behavioral data, learning directly from how a person communicates rather than from descriptions of their demographics or interview responses. This weight-level adaptation may overcome the “insufficient individuation” of prompting-based approaches.

105 This approach is grounded in the Interpersonal Theory of Suicide (ITS; Joiner 2005; Van Orden et al. 2010), which posits that the desire for suicide arises from the simultaneous presence of perceived burdensomeness and thwarted belonging—both inherently interpersonal constructs. If these constructs are reflected in communication patterns, they should manifest more
110 clearly in social interaction than in static text analysis.

The present study addresses six primary questions. First, do per-person fine-tuned language model agents produce risk-relevant language at rates that are associated with their real-world counterparts’ suicidal ideation? Second, does the arena provide incremental predictive validity over direct
115 analysis of participants’ actual text messages? Third, is the predictive signal specific to interpersonal communication style, or would any behavioral data suffice? Fourth, do standardized conversational probes—designed to target specific ITS constructs—differentiate from the unstructured all-pairs arena, and does construct-specific scoring improve prediction? Fifth, are
120 the adapters person-specific, such that clinical prediction depends on correct person–adapter matching? Sixth, does the learned communication style reflect stable individual differences or transient states?




Figure 1: Overview of the communicative stress test paradigm. Each participant’s real text messages are used to fine-tune a LoRA adapter on Qwen3-8B, creating a per-person agent. Each agent engages in 20-turn conversations with standardized probe personas. The generated text is scored with risk language dictionaries and correlated with EMA-measured suicidal ideation.

2. Method

2.1. Participants and Procedure

Participants were 79 adults (ages 18–65) recruited from the community who endorsed past-month active suicidal ideation or suicidal behaviors and owned an Android-based smartphone. Full eligibility criteria and recruitment procedures are described in [blinded for review]. Following a baseline assessment, participants completed 6 EMA prompts per day for 28 days (overall response rate = 68.8%). Simultaneously, a passive sensing application captured screenshots at five-second intervals during active smartphone use, yielding approximately 7.4 million screenshots across the sample ($M = 92,613$ per

participant).

2.2. Message Extraction

135 Text messages were extracted from screenshots using a vision-language model (VLM) pipeline. The Qwen2.5-VL model (Qwen Team, 2025) was deployed on NVIDIA L40S GPUs (48 GB VRAM). Each screenshot was resized to 400×672 pixels and processed with a single structured prompt containing eight questions (see Appendix Appendix A for full prompt text).
140 The prompt asked the VLM to identify the application being displayed, the type of activity, whether notifications were present, the nature of social media engagement, whether a keyboard was visible, and—if a keyboard was present in a messaging app—to extract text separately for the phone owner (typically right-aligned dialogue boxes) and the message recipient (typically
145 left-aligned). Generation used a maximum of 600 new tokens. Validation against 500 human-annotated screenshots demonstrated 99.7% accuracy for keyboard detection, and when keyboards were present, 100% accuracy for text extraction from both owner and recipient messages. This pipeline produced two text streams per participant—*owner text* (messages sent by the
150 participant) and *recipient text* (messages received).

2.3. EMA Measures

At each EMA prompt, participants reported on suicidal ideation using four items assessing passive ideation (wishes for death, thoughts of death) and active ideation (thoughts of killing oneself, intent to act). Additional
155 items assessed suicidal planning (3 items), desire for death (2 items), perceived burdensomeness (2 items), thwarted belonging (2 items), positive af-

fect (10 items from the PANAS), and negative affect (12 items). All items were summed within construct at each prompt, then aggregated to person-level means and standard deviations. We use both aggregates as outcome variables: means capture average symptom severity, while standard deviations capture within-person variability over the 28-day assessment period. Variability in suicidal ideation may be particularly clinically informative, as fluctuation in ideation has been identified as a predictor of the transition from ideation to suicidal behavior (Kleiman et al., 2017).

2.4. Per-Person Agent Training

2.4.1. Message Filtering

Because screenomics captures all on-screen activity, messages were extracted from any messaging application visible on the participant’s phone. The majority of extracted messages came from SMS/native messaging (86.7%), followed by social media messaging apps (13.1%), with smaller contributions from Snapchat, Discord, Messenger, WhatsApp, and others (see Appendix Appendix B for the full distribution). Participants used a median of 2 messaging platforms (range: 1–6). For each participant, sent (owner) text messages were cleaned in two stages. First, messages shorter than five characters, empty strings, and VLM output artifacts were removed. Artifact filtering targeted 16 patterns characteristic of VLM extraction failures. Second, a minimum message count threshold of 50 post-filtering messages was applied. Of the 79 participants, 73 met this threshold. Among eligible participants, message counts varied substantially ($Mdn = 733$; $M = 1,983$; range = 50–13,551). LoRA training was attempted for all 73; two failed to converge, yielding 71 agents with successful adapters. Of these, 64 could be

matched with valid EMA data.

2.4.2. LoRA Fine-Tuning

We used Qwen3-8B (Qwen Team, 2025) as the base language model. For
185 each participant, we trained a low-rank adaptation (LoRA; Hu et al. 2022)
adapter on that individual’s cleaned owner messages. LoRA modifies a small
number of parameters (rank $r = 8$, $\alpha = 16$, targeting the query and value
projection matrices) while keeping the base model frozen. Each adapter
was trained for 3 epochs with a causal language modeling objective. No
190 clinical labels were provided during training; the model learned only from
raw message text.

2.5. Arena Paradigms

We evaluated three conversational paradigms of increasing structure, each
using the same per-person LoRA adapters and identical generation param-
195 eters (temperature = 0.9, top- $p = 0.95$, repetition penalty = 1.3, maximum
200 new tokens per turn). All conversations were scored using the same
two-layer dictionary system described below.

2.5.1. Paradigm 1: All-Pairs Arena

Every possible pair of agents ($N = 2,485$ pairs from 71 agents) engaged
200 in a 20-turn simulated conversation seeded with “How are you feeling to-
day?” Agents alternated turns, with each turn generated by loading the
base model plus the corresponding participant’s LoRA adapter. A system
prompt established the conversational context: agents were instructed to re-
spond naturally and honestly, to keep responses to 1–3 sentences, and to

205 never repeat the previous message. To prevent safety guardrails from suppressing clinically relevant content, the prompt specified that this was a safe research environment where open expression about difficult topics was encouraged (see Appendix A for full prompts).

2.5.2. Paradigm 2: Caption-Trained Arena

210 To test whether the predictive signal originates specifically from interpersonal communication style, we conducted a comparison using an alternative training corpus. For each participant, a VLM (Qwen3-VL-2B-Instruct) had previously classified every screenshot (approximately 90,000 per participant) into structured behavioral categories using a JSON-output prompt (see Appendix Appendix A.1). Each screenshot was classified along six dimensions:
215 application category (16 options), activity type (10 options), content theme (9 options), mood tone (4 options), social context (6 options), and a brief visible text summary. These structured captions characterize the participant’s digital behavior patterns—but in the VLM’s categorical language rather than the participant’s own words. We fine-tuned the same Qwen3-8B
220 base model with an identical LoRA configuration on these caption corpora using one training epoch (reduced from three given the 50-fold increase in training data). Sixty-seven participants yielded successful caption-trained adapters. These agents were placed in the same all-pairs arena (2,211 paired
225 conversations, 20 turns each).

2.5.3. Paradigm 3: Standardized Probe Arena

To test whether structured conversational probes designed to elicit specific ITS constructs could outperform the unstructured all-pairs arena, we

developed five standardized probe personas, each targeting a distinct psychological dimension. Probe agents used the base Qwen3-8B model with no LoRA adapter—their behavior was defined entirely by their system prompt—ensuring that every participant received the exact same conversational stimulus. Each participant agent (base model + per-person LoRA) engaged in a 20-turn conversation with each of the five probes (71 participants $\times$ 5 probes = 355 conversations). The five probes were:

1. *Supportive Therapist* (target: emotional disclosure). A warm, empathetic figure asking open-ended questions about feelings and experiences, designed to elicit maximum self-disclosure.
2. *Friend in Crisis* (target: burdensomeness and help-seeking). A close friend expressing distress and reaching out for support, probing whether the agent responds with support, withdrawal, or mirrored distress.
3. *Rejecting Peer* (target: thwarted belonging). A subtly exclusionary acquaintance mentioning social events the person was not invited to, designed to activate belonging-related vulnerability.
4. *Demanding Authority* (target: perceived burdensomeness). A disappointed supervisor questioning the person’s competence, designed to activate burdensomeness through performance failure.
5. *Neutral Acquaintance* (target: baseline). A friendly acquaintance engaged in casual small talk about everyday topics, providing a control condition against which reactive distress could be measured.

In addition to the five original probes, a sixth probe—the Future Self—was included as an exploratory extension based on the temporal self-reflection

paradigm (Pataranutaporn et al., 2024). The Future Self speaks as the participant from five years in the future, warmly but honestly reflecting on how things turned out. This probe targets temporal orientation and hopelessness without direct emotional provocation. Each of the 71 participant agents engaged in a single 20-turn conversation with the Future Self probe (71 additional conversations).

Full probe persona prompts are provided in Appendix A. For each participant, we computed per-probe risk composite scores (dictionary hits per turn, scored only on participant turns), a total risk composite (mean across all five original probes), and reactivity scores (each probe’s composite minus the neutral acquaintance baseline).

2.6. Risk Language Scoring

To score conversation content, we employed a two-layer dictionary approach. The first layer consisted of the seven sub-dictionaries developed by Ammerman et al. (2026b), who adapted the 276-word crisis language dictionary validated by Swaminathan et al. (2023) (sensitivity = 0.98, PPV = 0.66). Two experts independently assigned each word to seven themes: suicidal thoughts (41 entries), non-substance methods (75 entries), substances (26 entries), sleep (7 entries), hopelessness (26 entries), general risk (93 entries), and help-seeking (8 entries).

The second layer added six supplementary categories aligned with the ITS (Van Orden et al., 2010), the Integrated Motivational-Volitional model (O’Connor, 2011), and absolutist thinking (Al-Mosaiwi & Johnstone, 2018): perceived burdensomeness (18 patterns), thwarted belonging (19 patterns), entrapment (14 patterns), negative affect (36 patterns), absolutist thinking

(18 patterns), and self-harm (5 patterns). The complete dictionary is provided in Supplemental Table S1.

280 2.7. Validation Analyses

Five analyses were conducted to assess the robustness and specificity of the clinical signal.

2.7.1. Adapter Shuffle Control

To test whether the clinical signal is person-specific—that is, whether
285 it resides in the learned adapter rather than in some artifact of the arena procedure—we randomly reassigned LoRA adapters to different participants (seed = 42) and re-ran the all-pairs arena with these mismatched adapters. If the signal is carried by person-specific adaptation, correlations between arena output and the *participant’s* SI should drop to zero when adapters are
290 mismatched. Conversely, if the adapter carries the signal, the shuffled arena output should correlate with the *adapter-owner’s* SI rather than the assigned participant’s.

2.7.2. Temporal Split Validation

To test whether adapters capture stable communicative dispositions rather
295 than transient states, we chronologically split each participant’s messages at the median timestamp, trained separate LoRA adapters on the first and second halves, and ran the full 5-probe battery with each set of adapters independently. We report: (a) split-half correlation of risk composites between the two temporal halves, (b) Spearman-Brown corrected reliability, and (c)
300 clinical prediction (correlation with EMA SI) for each split separately. Of

71 eligible participants, 66 had sufficient messages in each half (minimum 25 per half).

2.7.3. *Cross-Validated Prediction*

To assess out-of-sample predictive validity, we used 10-fold cross-validation
305 and leave-one-out cross-validation (LOOCV) with Ridge regression predicting SI from the five per-probe risk composites (5 features). Ridge regression was chosen to guard against overfitting given the small sample and moderate feature count.

2.7.4. *Prompt Ablation*

310 To test whether the clinical signal depends on the disclosure-encouraging language in the participant system prompt (a potential demand characteristic), we re-ran the full 5-probe battery with a minimal participant prompt that omitted all references to mental health, self-harm, and suicide. The minimal prompt read: “You are a person having a casual conversation with
315 someone you know. Respond naturally and honestly as yourself. Keep your response to 1-3 sentences. NEVER repeat or paraphrase what the other person just said. Always say something new and different from the previous message.” Probe persona prompts were unchanged. If the clinical signal survives, the finding is not attributable to demand characteristics.

320 2.7.5. *No-LoRA Floor Control*

To establish what proportion of risk language is attributable to the person adapter versus the base model, we ran the 5-probe battery 20 times using the base Qwen3-8B model with no LoRA adapter and the original

participant prompt. This establishes the floor: the base model’s intrinsic
325 tendency to generate risk-relevant content.

2.7.6. *Construct-Matched Feature Scoring*

To test whether construct-specific probe scoring improves prediction over
generic composites, we pre-specified a mapping from each probe to its stated
target construct (as defined in the probe design rationale, Section 2.5.3 and
330 Appendix A): Supportive Therapist → negative affect; Friend in Crisis →
perceived burdensomeness and help-seeking; Rejecting Peer → thwarted be-
longing; Demanding Authority → perceived burdensomeness. For the Neu-
tral Acquaintance (target: baseline), we used the overall risk composite. This
yielded six pre-specified features derived directly from the probe design rather
335 than from empirical performance.

2.8. *Sensitivity Analyses*

To address potential confounds, we conducted five sensitivity analyses:
(a) verbosity normalization—re-scoring all conversations with dictionary hits
normalized by token count rather than turn count; (b) absolutist thinking
340 specificity—testing whether absolutist language predicts SI above and be-
yond negative affect, replicating Al-Mosaiwi & Johnstone (2018); (c) training
heterogeneity—testing whether message count (range: 50–13,551) confounds
the risk composite–SI association; (d) reactivity regression to the mean—
testing whether negative reactivity correlations are mechanical artifacts by
345 computing semi-partial correlations controlling for neutral baseline; (e) se-
lection bias—comparing excluded participants ($N = 15$) to included partici-
pants ($N = 64$) on baseline SI and message volume.

2.9. Analytic Plan

All primary analyses were between-person, correlating each participant's arena-derived scores with their person-level EMA aggregates (both means and standard deviations). We report Pearson correlations for all primary analyses, with Spearman correlations for zero-inflated distributions. The central comparison examined three sources of risk language—real messages, all-pairs arena, and probe arena—against the same EMA outcomes. For the probe arena, we additionally tested construct-specific predictions (e.g., does the Rejecting Peer specifically predict thwarted belonging?) and reactivity scores (probe minus neutral baseline). The validation analyses described above assessed adapter specificity (shuffle control), temporal stability (split-half), out-of-sample generalization (cross-validation), demand characteristics (prompt ablation), and base-model floor (no-LoRA control). We applied no correction for multiple comparisons; this study is exploratory, and we report the total number of tests conducted to facilitate reader evaluation. We note that the present analyses do not incorporate baseline clinical data (e.g., baseline interview-assessed SI severity, diagnostic status, or demographic covariates) as predictors or covariates; an important question for future work is whether arena-derived risk scores provide incremental predictive validity above and beyond baseline assessment.

3. Results

Figure 2 provides a roadmap of all analyses reported in this section, organized by purpose: core paradigms testing the communicative stress test

concept, validation analyses confirming the signal’s properties, and sensitivity analyses addressing potential confounds.

3.1. Risk Language in Raw Messages

Before examining arena-generated language, we assessed the prevalence
375 of risk-relevant content in participants’ actual text messages. Of the 79
participants, 20 (25.3%) had at least one message containing explicit suicide-
related language (e.g., “suicidal,” “kill myself,” “end it all”), but such messages
were extremely rare—264 of 155,609 total messages (0.17%). Applying the
broader risk dictionary to raw messages, death/suicide terms appeared in
380 messages from 51% of participants, belonging-related terms in 38%, negative
affect in 39%, and hopelessness in 23%. Self-harm language was rare (5%
of participants). These base rates illustrate the core challenge that moti-
vates the arena approach: risk-relevant language is present in real messages
but sparse and diluted across thousands of routine communications, making
385 direct surveillance both ethically problematic and statistically inefficient.

3.2. All-Pairs Arena

3.2.1. Descriptives

The 71 agents produced 2,485 paired conversations (49,700 total turns)
with 0% echo rate. Of the 71 agents, 26 (37%) spontaneously mentioned
390 suicide at least once. After matching with EMA data, 64 participants consti-
tuted the analytic sample. Table 1 presents descriptive statistics for the EMA
outcome variables. Suicidal ideation scores were positively skewed ($Mdn =$
1.18, $M = 1.80$), with substantial between-person variability ($SD = 2.15$;

Table 1: Descriptive statistics for EMA outcome variables ($N = 64$).

Measure	M	SD	Mdn	Range
<i>Person-level means</i>				
Suicidal ideation (composite)	1.80	2.15	1.18	0.00–8.50
Passive SI	2.98	1.21	2.57	2.00–7.00
Active SI	2.82	1.04	2.42	2.00–6.22
Suicidal planning	3.25	0.53	3.09	2.98–6.67
Desire for death	2.49	0.78	2.18	2.00–5.33
Perceived burdensomeness	2.35	2.28	1.66	0.00–7.67
Thwarted belonging	2.72	2.23	2.40	0.00–7.92
Negative affect (sum)	7.78	6.61	6.40	0.10–30.60
Positive affect (sum)	9.44	6.75	7.90	0.98–29.97
<i>Within-person variability (person-level SDs)</i>				
Passive SI	0.93	0.77	0.87	0.00–3.30
Active SI	0.85	0.65	0.81	0.00–2.51

range: 0.00–8.50). Within-person variability in SI was also substantial: pas-
395 sive SI SD averaged 0.93 and active SI SD averaged 0.85 across participants,
indicating meaningful fluctuation over the 28-day assessment period.

3.2.2. Suicide Mention Rate Is Associated with Real Suicidal Ideation

The rate at which agents mentioned suicide was significantly associated
with real-world SI (Pearson $r = .269$, $p = .032$; Spearman $\rho = .438$, $p <$
400 $.001$). Agents who mentioned suicide at least once had significantly higher
real SI ($M = 2.78$) than those who never mentioned it ($M = 1.22$), $t(62) =$

2.98, $p = .004$, Cohen's $d = 0.77$.

3.2.3. Risk Dictionary Correlations

The strongest correlate of mean SI was absolutist thinking ($r = .607$,
405 $p < .001$), followed by thwarted belonging ($r = .474$, $p < .001$), perceived
burdensomeness ($r = .403$, $p = .001$), substance use ($r = .375$, $p = .002$),
and death ideation ($r = .349$, $p = .005$). The full risk composite correlated at
 $r = .433$ ($p < .001$). Absolutist thinking was the only category significantly
associated with every clinical construct assessed, including suicidal planning
410 ($r = .416$) and desire ($r = .544$).

3.2.4. Construct Convergence

Simulated burdensomeness was associated with real perceived burden-
someness ($r = .282$, $p = .024$); simulated thwarted belonging was associated
with real thwarted belonging ($r = .323$, $p = .009$). Simulated negative affect
415 was not associated with EMA negative affect ($r = .194$, $p = .124$), suggesting
the ITS interpersonal constructs are better captured than general emotional
distress.

3.2.5. Arena Versus Raw Messages

Across 112 head-to-head comparisons, the arena won 66%, real sent mes-
420 sages won 19%, and sent + received messages won 15%. The arena's advan-
tage was concentrated in interpersonal constructs: burdensomeness (owner
.079, arena .403), thwarted belonging (owner .217, arena .474), absolutist
thinking (owner .113, arena .607). Raw messages outperformed for explicit
crisis content: non-substance methods (owner .450, arena .073) and hope-
425 lessness (owner .399, arena .306).

3.3. *Caption-Trained Arena Produces No Clinical Signal*

The caption-trained arena yielded no significant associations with SI. The risk composite was not associated with mean SI ($r = .159$, $p = .198$), compared to $r = .433$ ($p < .001$) for the message-trained arena. No individual
430 dictionary category reached significance. Absolutist thinking ($r = -.044$, $p = .724$), thwarted belonging ($r = .077$, $p = .536$), and perceived burdensomeness ($r = .062$, $p = .617$) were all near zero. The caption-trained arena won only 25% of 112 comparisons, versus 41% for real messages and 34% for combined messages. Despite having approximately $50\times$ more training data,
435 the caption-trained agents failed to capture the clinically relevant signal that emerged when agents were trained on participants' own words.

3.4. *Standardized Probe Arena*

3.4.1. *Descriptives*

The 71 participant agents engaged in 355 conversations (5 probes each,
440 20 turns). Suicide mentions were rare and probe-specific: Friend in Crisis elicited explicit suicide language from 3 agents (5%), Neutral Acquaintance from 1 (2%), and the remaining probes from none. Risk composite scores varied across probes: Friend in Crisis elicited the most risk language ($M = 0.52$, $SD = 0.59$), followed by Demanding Authority ($M = 0.45$, $SD =$
445 0.54), Supportive Therapist ($M = 0.39$, $SD = 0.58$), Neutral Acquaintance ($M = 0.35$, $SD = 0.52$), and Rejecting Peer ($M = 0.33$, $SD = 0.43$). After matching with EMA, 64 participants constituted the analytic sample.

3.4.2. Probe Battery Prediction

The mean risk composite across all five probes was associated with SI at
450 $r = .576$ ($p < .001$)—stronger than the all-pairs arena ($r = .477$), real sent
messages ($r = .461$), or any single analysis approach. This advantage was
consistent across nearly all EMA outcomes: passive SI (probes .585, arena
.470, owner .454), active SI (probes .514, arena .442, owner .425), planning
(probes .334, arena .300, owner .220), desire (probes .357, arena .278, owner
455 .238), and negative affect (probes .375, arena .244, owner .312). Real mes-
sages outperformed only for perceived burdensomeness (.440 vs. .392) and
thwarted belonging (.441 vs. .368). However, formal comparison of dependent
correlations (Steiger’s test and bootstrap with 10,000 resamples) confirmed
that none of the pairwise differences among probes reached significance at
460 $N = 64$, with the exception of Neutral Acquaintance versus Demanding Au-
thority ($z = 3.34$, $p < .001$; bootstrap 95% CI for difference: [.119, .704]).
The combined probe composite ($r = .576$) did not significantly differ from
the Neutral Acquaintance alone ($r = .551$; $z = 0.29$, $p = .774$; bootstrap
95% CI: [−.217, .314]) or from the Supportive Therapist ($r = .412$; $z = 1.72$,
465 $p = .085$). Cross-validation further reveals that a two-probe combination
outperforms the full battery (see Section 3.5).

3.4.3. The Neutral Acquaintance Is the Best Single Probe

Among individual probes, the Neutral Acquaintance—a casual small-talk
partner with no emotional provocation—showed the single strongest associ-
470 ation with SI ($r = .551$, $p < .001$), outperforming the Supportive Therapist
($r = .412$, $p < .001$), Rejecting Peer ($r = .326$, $p = .009$), Friend in Crisis
($r = .316$, $p = .011$), and Demanding Authority ($r = .119$, $p = .348$). This

Table 2: Correlations between probe-derived risk composites and EMA outcomes ($N = 64$). Bold indicates $p < .01$; * indicates $p < .05$.

Source	SI	Passive	Active						Neg
	Mean	SI	SI	Plan	Desire	Burden	Belong	Aff	
Neutral Acquaintance	.551	.557	.493	.189	.492	.259*	.271*	.314*	
Supportive Therapist	.412	.416	.368	.168	.270*	.222	.249*	.309	
Future Self	.369	—	—	—	—	.191	.253*	.262*	
Rejecting Peer	.326	.290*	.293*	.249*	.112	.374	.304	.119	
Friend in Crisis	.316*	.277*	.276*	.227	.200	.263*	.250*	.173	
Demanding Authority	.119	.144	.098	.009	.023	.080	.038	.131	
Combined (5 probes)	.576	.585	.514	.334	.357	.392	.368	.375	
All-pairs arena	.477	.470	.442	.300	.278*	.267*	.323*	.244	
Raw messages (owner)	.461	.454	.425	.220	.238	.440	.441	.312*	

pattern held for passive SI (neutral .557, therapist .416), active SI (neutral .493, therapist .368), and desire (neutral .492, therapist .270). The Neutral
475 Acquaintance was also significantly associated with perceived burdensomeness ($r = .259$, $p = .039$), thwarted belonging ($r = .271$, $p = .030$), and negative affect ($r = .314$, $p = .012$). Figure 4 displays the per-probe correlations; Table 2 presents the full omnibus correlation matrix across all probes and EMA outcomes.

480 3.4.4. Construct-Specific Predictions

The Rejecting Peer’s risk composite was significantly associated with its target construct, real thwarted belonging ($r = .304$, $p = .015$), and with SI ($r = .326$, $p = .009$). The Friend in Crisis was associated with its target, perceived burdensomeness ($r = .263$, $p = .036$), and SI ($r = .316$, $p =$

485 .011). The Demanding Authority failed entirely: it was not associated with
perceived burdensomeness ($r = .080$, $p = .530$) or SI ($r = .119$, $p = .348$).
However, these construct-specific probe correlations did not exceed the all-
pairs arena's corresponding values (arena thwarted belonging: $r = .323$;
arena burdensomeness: $r = .267$), indicating that the targeted probes did
490 not achieve superior construct-specific prediction.

3.4.5. Reactivity Scores Do Not Predict SI

Reactivity scores—computed as each probe's risk composite minus the
neutral acquaintance baseline—were uniformly not associated with SI. The
Rejecting Peer reactivity showed a marginal negative association with SI
495 ($r = -.232$, $p = .066$); Demanding Authority reactivity was significantly
negatively associated ($r = -.350$, $p = .005$). These negative associations in-
dicate that participants with higher SI produced *less* additional risk language
in response to interpersonal provocation relative to their already-elevated
baseline, consistent with the interpretation that the signal resides in tonic
500 communication style rather than reactive distress.

3.5. Adapter Specificity: The Signal Is Person-Specific

When adapters were shuffled—randomly reassigned to different participants—
the risk composite correlation with the *original participant's* SI dropped to
 $r = .071$ ($p = .577$), compared to $r = .433$ ($p < .001$) for correctly matched
505 adapters. Critically, the shuffled adapter's output correlated $r = .426$ ($p < .001$)
with the *adapter-owner's* SI—the person whose messages trained the
adapter (Figure 5). This double dissociation confirms that each adapter en-
codes its own person's clinical signal: the signal follows the adapter, not the

participant to whom it is assigned. When a high-risk person’s adapter is
510 loaded onto a different participant’s slot, the generated language still reflects
the high-risk person’s communicative style.

3.6. *Construct-Matched Scoring Does Not Improve Prediction*

The pre-specified construct-matched features—scoring each probe only on
its stated target construct (e.g., Rejecting Peer scored only on thwarted be-
515 longing, Friend in Crisis scored only on perceived burdensomeness)—did not
improve prediction over simple probe composites. The six construct-matched
features achieved LOOCV $r = .358$ ($p = .004$), compared to LOOCV $r =$
.476 for the Neutral Acquaintance composite alone and LOOCV $r = .462$
for the five probe composites. Without the Neutral Acquaintance base-
520 line, the four construct-specific features alone yielded LOOCV $r = .106$
($p = .405$), failing entirely. Individual construct-matched correlations with
SI were mixed: Friend in Crisis $\rightarrow$ perceived burdensomeness ($r = .320$,
 $p = .010$) and Rejecting Peer $\rightarrow$ thwarted belonging ($r = .335$, $p = .007$)
reached significance, but Supportive Therapist $\rightarrow$ negative affect ($r = -.017$,
525 $p = .894$) and Demanding Authority $\rightarrow$ perceived burdensomeness ($r = .109$,
 $p = .392$) did not. These results indicate that the probes designed to target
specific ITS constructs do not reliably elicit construct-specific language at
levels sufficient for differential prediction.

3.7. *Cross-Validated Probe Battery Optimization*

530 The Future Self probe was associated with SI at $r = .369$ ($p = .003$),
positioned between the Neutral Acquaintance ($r = .551$) as the strongest
individual probe and the Demanding Authority ($r = .119$) as the weakest.

Exhaustive subset search identified {Neutral Acquaintance, Supportive Therapist, Future Self} as the optimal 3-probe combination in-sample ($r = .655$);
535 however, nested cross-validation—performing subset selection independently within each held-out fold—reduced this to LOOCV $r = .420$ ($p < .001$) and 10-fold $r = .492$ ($p < .001$), confirming substantial overfitting in the in-sample estimate. The 3-probe subset was selected in 83% of LOOCV folds, indicating stable feature selection despite attenuated effect size.

540 Importantly, the simpler fixed combination of Neutral Acquaintance and Supportive Therapist—requiring no data-driven selection—achieved LOOCV $r = .570$ ($p < .001$), outperforming both the nested-CV optimized subset ($r = .420$) and the full 5-probe battery (LOOCV $r = .462$). The addition of any further probes reduced cross-validated performance (Figure 6). This
545 suggests that two probes providing complementary conversational contexts—one allowing free self-expression, the other providing structured emotional elicitation—capture the clinical signal more efficiently than a larger battery that introduces situation-driven noise.

3.8. Temporal Stability

550 First-half and second-half risk composites correlated at $r = .368$ ($p = .002$; Spearman-Brown corrected reliability = .538). Both splits were significantly associated with SI: first-half $r = .359$ ($p = .005$), second-half $r = .396$ ($p = .002$). However, construct-specific predictions diverged: second-half adapters recovered perceived burdensomeness ($r = .34$) and thwarted be-
555 longing ($r = .40$) correlations comparable to full adapters, while first-half adapters did not ($r = .08-.12$, n.s.). The Neutral Acquaintance was the most temporally stable probe (first vs. second half: $r = .304$, $p = .013$). For

comparison, EMA suicidal ideation scores themselves showed split-half reliability of $r = .451$ (Spearman-Brown = .621; $N = 49$), with burdensomeness
560 at $r = .668$ (SB = .801) and belonging at $r = .722$ (SB = .839). The adapter split-half ($r = .368$, SB = .538) is lower than the EMA SI benchmark but within the same order of magnitude, indicating that the adapters capture a substantial portion—though not all—of the temporal stability present in the clinical outcome itself. The gap between adapter and EMA stability
565 likely reflects both measurement noise from the dictionary-based scoring and genuine state-dependent variation in communication style (Figure 7).

3.9. Cross-Validated Prediction

Using the five original per-probe risk composites as features, Ridge regression with LOOCV yielded $r = .462$ ($p < .001$). The two-probe configura-
570 tion (Neutral Acquaintance and Supportive Therapist) achieved the highest cross-validated prediction at LOOCV $r = .570$ ($p < .001$), outperforming the full battery. A single probe (Neutral Acquaintance alone) achieved LOOCV $r = .476$ ($p < .001$). These results confirm that out-of-sample prediction is genuine but that additional probes beyond two do not improve general-
575 ization. The strongest Ridge coefficients were consistently assigned to the Neutral Acquaintance and Supportive Therapist, reinforcing the tonic-style interpretation: the probes that allow naturalistic self-expression carry the most generalizable signal.

3.10. Prompt Ablation: The Signal Survives Without Disclosure Priming

580 With the minimal prompt, the risk composite was associated with SI at $r = .430$ ($p < .001$), compared to $r = .576$ ($p < .001$) with the original

prompt. Associations with perceived burdensomeness ($r = .309$, $p = .013$) and thwarted belonging ($r = .289$, $p = .021$) also survived. Only negative affect dropped to non-significance ($r = .191$, $p = .130$). Cross-prompt com-
585 posites correlated at $r = .612$ ($p < .001$), confirming that the same individual differences are expressed regardless of prompt condition. The prompt’s role is therefore one of permission—amplifying an existing signal rather than creating one.

The per-probe pattern shifted under the minimal prompt in a theoretically
590 cally informative way. The Rejecting Peer became the strongest single predictor (Spearman $\rho = .468$, $p < .001$; Pearson $r = .410$ after outlier removal), while the Neutral Acquaintance dropped from $r = .551$ to $r = .235$ ($p = .061$). This inversion suggests that without explicit disclosure permission, the clinical signal concentrates in contexts that naturally create
595 interpersonal pressure, whereas with disclosure permission, it pervades even casual conversation. Notably, suicide mentions did not decrease under the minimal prompt (2.5% vs. 1.1%), counter to the demand-characteristic prediction (Figure 8).

3.11. *No-LoRA Floor Control*

600 The base model without any LoRA adapter produced risk language at a mean composite rate of 0.0036 ($SD = 0.0011$), nearly identical to the LoRA-adapted mean of 0.0035 ($SD = 0.0017$). Zero of 100 base-model conversations contained suicide mentions. The critical difference was variance: LoRA-adapted agents showed 2.44 times the between-conversation variance
605 of the base model. The LoRA’s contribution is not in shifting the mean level of risk language but in creating individual-level differentiation that tracks

real clinical differences (Figure 9).

3.12. Sensitivity Analyses

3.12.1. Verbosity

610 Mean tokens per turn correlated modestly with SI ($r = .384$, $p = .002$), confirming a potential confound. However, token-normalized risk composites were still significantly associated with SI ($r = .438$, $p < .001$), with only modest attenuation from the per-turn rate ($r = .576$). Suicidal thoughts, substance use, thwarted belonging, and general risk categories survived token
615 normalization.

3.12.2. Absolutist Thinking Specificity

Absolutist thinking and negative affect were nearly uncorrelated in agent-generated text ($r = .057$), unlike in human-authored samples where they moderately overlap. When controlling for negative affect, the partial corre-
620 lation of absolutist thinking with SI was $r = .245$ ($p = .051$), which was no longer significant, highlighting the shared variance between negative affect and absolutist language in this paradigm.

3.12.3. Training Heterogeneity

Message count correlated weakly with SI ($r = .224$, $p = .075$) and mod-
625 estly with risk composite ($r = .310$, $p = .013$). After partialing out message count, the risk composite was still strongly associated with SI (partial $r = .547$, $p < .001$). However, a median split revealed that prediction was substantially stronger in the high-volume half ($r = .610$, $p < .001$) than the low-volume half ($r = .186$, n.s.), indicating that adapters trained on
630 more data produce more valid signals. The practical implication is that a

minimum message volume is needed for reliable assessment. A systematic titration study—retraining adapters on progressively smaller message subsets (e.g., 50, 100, 200, 500, 1000) and measuring the degradation in clinical association—would establish the precise minimum, but is beyond the scope
635 of the present work.

3.12.4. *Reactivity Regression to the Mean*

The negative association between reactivity scores and SI was not a mechanical artifact: reactivity was associated with SI above and beyond the neutral baseline in multiple regression ($\beta = 2.081$, $p = .003$; model $R^2 = .397$),
640 and the semi-partial correlation of reactivity residuals with SI was $r = .306$ ($p = .014$).

3.12.5. *Selection Bias*

The 15 participants excluded from the analytic sample had somewhat higher SI than the 64 included, though this difference was not statistically
645 significant when computed on the same scale ($M = 2.20$ vs. $M = 1.78$; $U = 123$, $p = .150$; $N_{\text{excluded with EMA}} = 6$). The excluded group had dramatically fewer messages ($M = 12$ vs. $M = 1,202$; $p < .001$). While the SI difference is modest, the message volume difference is extreme—the method cannot model individuals who produce minimal digital text, regardless of their risk
650 level.

4. Discussion

This study introduces a communicative stress test paradigm for studying suicidal vulnerability—training per-person language model agents on real

text messages and placing them in simulated social interactions. Several
655 principal findings emerge, each with distinct theoretical and practical implications.

4.1. Communication Style Encodes Suicidal Vulnerability

Per-person fine-tuned agents produce risk-relevant language at rates that
meaningfully predict real-world suicidal ideation, despite never being trained
660 on clinical data. The strongest single correlate—absolutist language ($r = .607$)—replicates and extends the seminal finding of Al-Mosaiwi & Johnstone (2018) in a novel paradigm. The ITS-grounded construct convergence provides preliminary evidence that the arena captures theoretically meaningful interpersonal processes: simulated burdensomeness predicted real burdensomeness; simulated belonging predicted real belonging.
665

4.2. The Signal Is Specifically Interpersonal

The caption-trained arena comparison provides a critical specificity test. Agents trained on VLM-generated descriptions of participants’ entire digital lives—approximately 90,000 screenshots per person covering all applications, content, and temporal patterns—produced no clinical signal whatsoever. This rules out two alternative explanations: that any behavioral data
670 would work (it does not), and that more training data is better ($50\times$ more data produced worse results). The predictive signal resides specifically in *how a person talks to other people*, not in what they do on their phone.

675 4.3. Tonic Style, Not Probe Engineering, Drives Prediction

The standardized probe arena ($r = .576$ for combined probes predicting SI) outperformed both the all-pairs arena ($r = .477$) and raw message

analysis ($r = .461$). In-sample, the five-probe composite ($r = .576$) and the single Neutral Acquaintance ($r = .551$) appear nearly equivalent. However, cross-validation reveals a more nuanced picture: the Neutral Acquaintance alone achieves LOOCV $r = .476$, while adding the Supportive Therapist yields LOOCV $r = .570$ —a genuine incremental gain that is obscured in the in-sample comparison. Adding further probes beyond these two reduces out-of-sample prediction (full 5-probe LOOCV $r = .462$), and the in-sample optimal 3-probe subset ($r = .655$) was substantially overfit, dropping to $r = .420$ under nested cross-validation—a cautionary example of post-hoc probe selection on small samples.

Theory-driven probe engineering did not achieve its intended purpose of construct-specific prediction. Each probe was designed to target a specific ITS construct (e.g., Rejecting Peer for thwarted belonging, Demanding Authority for perceived burdensomeness), and the dictionaries contain categories aligned with these constructs. However, scoring each probe only on its target construct did not improve prediction over simple composites: the strictly pre-specified construct-matched features (LOOCV $r = .358$) performed worse than the Neutral Acquaintance composite alone (LOOCV $r = .476$). The Rejecting Peer predicted thwarted belonging ($r = .304$) but not better than the all-pairs arena ($r = .323$); the Demanding Authority failed entirely. The probes do not manipulate what they claim to manipulate at a level that translates into differential prediction.

Instead, what the agents encode is tonic communication style—a person-level disposition that pervades all conversational contexts rather than a construct-specific reactivity that appears only under targeted provocation. The nega-

tive reactivity correlations confirm this: higher-SI participants produced *less* additional risk language when provoked, not more, suggesting their baseline
705 is already elevated and further provocation pushes them toward withdrawal rather than expression.

The most striking finding is that the Neutral Acquaintance—casual small talk with no emotional provocation—was the single best predictor of SI ($r = .551$). This result has a clear theoretical interpretation grounded in
710 the Fluid Vulnerability Theory (Bryan & Rudd, 2018): chronic suicidal vulnerability pervades everyday communication at a level that is detectable even in conversations about weather and weekend plans. The provocative probes do not add construct-specific signal; their value, when it exists (as with the Supportive Therapist), lies in providing a complementary conversational con-
715 text that elicits a different facet of the same tonic style. Why the Supportive Therapist specifically adds incremental value over the Neutral Acquaintance is not clear from the present data. Possible mechanisms include encouraging emotional elaboration (increasing the density of scorable content), providing a complementary conversational register, or eliciting a different facet of tonic
720 style. Future work should investigate this directly, for example by comparing token counts and content diversity across probe conditions.

4.4. *The Disclosure Prompt as Permission Gate*

The prompt ablation demonstrates that the clinical signal partially survives the removal of disclosure-encouraging language. With the minimal
725 prompt—containing no references to mental health, self-harm, or suicide—the risk composite was still associated with SI at $r = .430$ ($p < .001$), a 26% relative reduction from the original $r = .576$. Cross-prompt composites

correlated at $r = .612$ ($R^2 = .375$), indicating that approximately 38% of the between-person variance in risk language is shared across prompt conditions, with the remaining 62% prompt-dependent. The disclosure language substantially amplifies the signal but does not create it. This is more accurately described as “demand characteristics account for some but not all of the effect” rather than “demand characteristics are ruled out.”

Instead, the disclosure prompt functions as a permission gate—a contextual variable that determines *where* vulnerability is expressed, not *whether* it exists. With disclosure permission, tonic vulnerability pervades all conversational contexts including neutral small talk; without it, the same vulnerability concentrates in contexts that naturally create emotional pressure (social rejection). The Rejecting Peer’s rise from $r = .326$ to Spearman $\rho = .468$ under the minimal prompt, while the Neutral Acquaintance dropped from $r = .551$ to $r = .235$, demonstrates this redistribution clearly. This parallels findings in clinical interviewing, where structured permission to discuss suicidal thoughts increases disclosure without altering underlying risk (Joiner, 2005).

Notably, suicide mentions did not decrease under the minimal prompt (2.5% vs. 1.1%), directly contradicting the demand-characteristic hypothesis. If the disclosure-encouraging language were manufacturing risk content, its removal should have reduced suicide mentions; instead, they increased, suggesting the minimal prompt may have reduced compliance with conversational norms that typically suppress such content.

4.5. *The Floor Beneath the Signal*

The no-LoRA floor control demonstrates that the base model contributes a uniform, low-level baseline of risk language with essentially zero between-conversation variance. The mean risk composite was nearly identical with
755 and without LoRA (0.0035 vs. 0.0036), but the LoRA-adapted agents showed 2.44 times the variance. The adapter’s contribution is not to increase risk language overall but to differentiate individuals—creating a distribution of risk language rates where the rank ordering maps onto real clinical differences.

This finding distinguishes our approach from concerns about LLM “hal-
760 lucination” of clinical content. The risk language is not model-generated noise uniformly distributed across all agents; it is person-specific variation introduced by the LoRA adapter and traceable to individual communicative style. Combined with the shuffle validation (Section 3.4), this provides converging evidence that the clinical signal is genuinely person-level: it follows
765 the adapter, not the base model, and not the participant slot.

4.6. *Adapter Specificity and the Digital Twin Challenge*

The shuffle validation provides the strongest evidence that the clinical signal is person-specific. When adapters were mismatched, the signal followed the adapter ($r = .426$ with the adapter-owner’s SI) rather than disappearing
770 into noise—demonstrating that LoRA fine-tuning creates genuine individual-level representations. The double dissociation (prediction drops to $r = .071$ for the wrong participant, remains at $r = .426$ for the adapter-owner) rules out the possibility that the clinical signal is an artifact of arena position, conversation dynamics, or dictionary scoring.

775 This result contrasts sharply with the performance of prompt-based digital twins. Peng et al. (2025) identified five systematic distortions in LLM-based digital twins and found that the average correlation between simulated and real human responses is only $r = 0.20$. Our LoRA-based approach may overcome the “insufficient individuation” problem precisely because it learns
780 from actual behavioral data rather than from demographic or interview-based descriptions. Whereas prompt-based conditioning provides the model with *information about* a person, weight-based adaptation gives the model *experience as* that person—a fundamentally different personalization mechanism that appears to encode deeper aspects of communicative identity.

785 4.7. Temporal Stability and the Nature of the Learned Signal

The split-half reliability of .538 falls below the conventional .70 threshold for stable individual differences, suggesting that the adapters capture a mixture of stable communicative dispositions and time-varying patterns. This is not necessarily a limitation. The Fluid Vulnerability Theory (Bryan & Rudd,
790 2018) posits that suicidal vulnerability fluctuates; a method that captures both stable and dynamic components may be more clinically informative than one that captures only trait-level differences. Crucially, SI prediction was preserved across both temporal splits ($r = .359$ and $r = .396$), indicating that the core clinical signal is sufficiently stable for practical use.

795 The finding that second-half adapters recovered construct-specific predictions (burdensomeness $r = .34$, belonging $r = .40$) better than first-half adapters ($r = .08$ – $.12$, n.s.) may reflect two non-mutually-exclusive processes: participants’ increasing comfort with texting over the study period, producing more naturalistic and self-revealing messages in later weeks; and

800 temporal proximity between later messages and EMA assessments, resulting in better alignment between communication style and concurrent clinical state. The Neutral Acquaintance’s status as the most temporally stable probe ($r = .304$) is consistent with its role as a measure of tonic communication style: what a person reveals in casual small talk changes less over time than
805 what they reveal under specific interpersonal pressure.

4.8. Who This Method Reaches—and Who It Misses

The 15 participants excluded from the analytic sample had dramatically fewer messages ($M = 12$ vs. $M = 1,202$; $p < .001$) and somewhat higher SI, though this difference did not reach significance ($M = 2.20$ vs. $M = 1.78$;
810 $p = .150$; only 6 excluded participants had EMA data). This reveals a fundamental boundary condition of any text-based approach: those who communicate least digitally are hardest to model. The method is best understood as a screening tool for moderate-risk populations who produce sufficient digital text, rather than as a universal risk assessment instrument. Training hetero-
815 geneity further constrains the method’s reach: prediction was substantially stronger in the high-message-volume half of the sample ($r = .610$) than the low-volume half ($r = .186$, n.s.), suggesting a practical minimum of approximately 700 messages for reliable clinical prediction.

4.9. Practical Implications

820 These findings suggest a deployment framework with two key parameters: message volume and probe configuration. Preliminary evidence from a median split suggests that adapters trained on more messages produce stronger associations (high-volume $r = .610$ vs. low-volume $r = .186$), though the

precise minimum message count for reliable assessment requires further systematic testing. Once trained, a single 20-turn conversation with a neutral small-talk partner provides a minimum viable risk estimate (LOOCV $r = .476$) that outperforms direct analysis of thousands of real messages. Adding a Supportive Therapist conversation provides the most robust cross-validated performance (LOOCV $r = .570$), while larger probe batteries do not improve generalization.

The pipeline operates on learned communication style rather than surveillance of ongoing private messages. Training the per-person adapter requires access to an individual’s text messages, but once trained, ongoing risk monitoring operates on generated text. This is an incremental rather than categorical improvement in privacy—reducing but not eliminating the need for access to private communications. The ethical implications of creating person-specific language models from private communications require careful consideration, even when those models are used only for research purposes.

4.10. Limitations and Future Directions

The sample is moderate ($N = 64$) and recruited for recent suicidal ideation. All analyses are between-person and correlational; the strongest next step would examine whether changes in arena behavior across sessions track within-person SI fluctuations. The association between absolutist thinking and SI was no longer significant after controlling for negative affect (partial $r = .245$, $p = .051$), suggesting shared variance between these constructs in agent-generated text. Token normalization attenuated some effects (risk composite from $r = .576$ to $r = .438$), confirming that verbosity partially confounds per-turn scoring, though core findings survive.

The dictionary-based scoring is context-insensitive (“I don’t want to kill
850 myself” scores identically to “I want to kill myself”); averaging over many
turns mitigates but does not eliminate this limitation. The two-probe bat-
tery recommendation (LOOCV $r = .570$) requires validation in an indepen-
dent sample, as does the in-sample optimal 3-probe subset ($r = .655$), which
dropped to $r = .420$ under nested cross-validation. No correction for multi-
855 ple comparisons was applied; this study is explicitly exploratory, generating
hypotheses for confirmatory testing. The validation challenges highlighted
by Larooij & Törnberg (2025b) remain relevant: agent fidelity sufficient for
clinical decision-making has not been established.

4.11. *Conclusions*

860 Suicidal vulnerability is encoded in everyday communication patterns—
in the words people choose, the cognitive rigidity of their thinking, and the
interpersonal themes that pervade their messages. These signals are diluted
in routine digital interactions but can be surfaced through simulated social
interaction. The clinical signal is demonstrably person-specific (following
865 the adapter in shuffle validation), partially survives removal of disclosure-
encouraging language, and reflects individual differentiation introduced by
the LoRA adapter against a uniform base-model floor. Tonic communication
style—what emerges even in casual conversation—carries more predictive
information than construct-specific reactivity. A two-probe battery (Neutral
870 Acquaintance and Supportive Therapist) achieves the best cross-validated
prediction (LOOCV $r = .570$), demonstrating that conversational diversity
outweighs targeted provocation.

References

- Aher, G. V., Arriaga, R. I., & Kalai, A. T. (2023). Using large language
875 models to simulate multiple humans and replicate human subject studies.
In *Proceedings of ICML 2023*.
- Al-Mosaiwi, M., & Johnstone, T. (2018). In an absolute state: Elevated use
of absolutist words is a marker specific to anxiety, depression, and suicidal
ideation. *Clinical Psychological Science*, 6(4), 529–542.
- 880 [Blinded for review]. (2026a). Beyond screen time: How details of smartphone
interactions map onto changes in suicidal ideation. Manuscript submitted
for publication.
- [Blinded for review]. (2026b). Smartphone-based text obtained via passive
sensing as it relates to direct suicide risk assessment. Manuscript submitted
885 for publication.
- Auerbach, R. P., et al. (2023). Intensive longitudinal assessment of adoles-
cents to predict suicidal thoughts and behaviors. *Journal of the American
Academy of Child & Adolescent Psychiatry*, 62(9), 1010–1020.
- Auerbach, R. P., et al. (2025). Using smartphone GPS data to detect the
890 risk of adolescent suicidal thoughts and behaviors. *JAMA Network Open*,
8(1), e2454247.
- Argyle, L. P., et al. (2023). Out of one, many: Using language models to
simulate human samples. *Political Analysis*, 31(3), 337–351.

- Belsher, B. E., et al. (2019). Prediction models for suicide attempts and
895 deaths: A systematic review and simulation. *JAMA Psychiatry*, 76(6),
642–651.
- Brickman, J., Gupta, M., & Oltmanns, J. R. (2025). Large language mod-
els for psychological assessment: A comprehensive overview. *Advances in
Methods and Practices in Psychological Science*.
- 900 Bryan, C. J., & Rudd, M. D. (2018). Nonlinear change processes and the
emergence of suicidal behavior. *New Ideas in Psychology*, 51, 34–37.
- Franklin, J. C., et al. (2017). Risk factors for suicidal thoughts and behaviors:
A meta-analysis of 50 years of research. *Psychological Bulletin*, 143(2),
187–232.
- 905 Holmes, E. A., et al. (2025). Applications of large language models in suicide
prevention: Scoping review. *Journal of Medical Internet Research*, 27(1),
e63126.
- Hu, E. J., et al. (2022). LoRA: Low-rank adaptation of large language models.
In *Proceedings of ICLR*.
- 910 Huang, Y.-H., et al. (2024). Mindsets of suicide trajectories: An LIWC
analysis of hotline conversations. *Suicide and Life-Threatening Behavior*,
54(6), 1088–1101.
- Joiner, T. E. (2005). *Why people die by suicide*. Harvard University Press.
- Kambeitz, J., & Meyer-Lindenberg, A. (2025). Modelling the impact of

- 915 environmental and social determinants on mental health using generative
agents. *npj Digital Medicine*, 8, Article 36.
- Kim, Y., et al. (2025). CoDaS: AI co-data-scientist for biomarker discovery
via wearable sensors. *arXiv preprint arXiv:2604.14615*.
- Kleiman, E. M., et al. (2017). Examination of real-time fluctuations in suici-
920 dal ideation and its risk factors. *Journal of Abnormal Psychology*, 126(6),
726–738.
- Large, M., et al. (2017). Meta-analysis of longitudinal cohort studies of
suicide risk assessment. *PLoS ONE*, 12(6), e0180292.
- Li, R., et al. (2025). How far are LLMs from being our digital twins? A
925 benchmark for persona-based behavior chain simulation. In *Findings of
ACL 2025*.
- Louie, R., et al. (2024). Roleplay-doh: Enabling domain-experts to create
LLM-simulated patients via eliciting and adhering to principles. In *Pro-
ceedings of EMNLP 2024*.
- 930 O’Connor, R. C. (2011). Towards an integrated motivational-volitional model
of suicidal behaviour. In R. C. O’Connor, S. Platt, & J. Gordon (Eds.),
International handbook of suicide prevention (pp. 181–198). Wiley.
- Ophir, Y., et al. (2022). Linguistic features of suicidal thoughts and behav-
iors: A systematic review. *Clinical Psychology Review*, 95, 102161.
- 935 Park, J. S., et al. (2023). Generative agents: Interactive simulacra of human
behavior. In *Proceedings of the 36th ACM UIST*.

- Park, J. S., et al. (2024). Generative agent simulations of 1,000 people. *arXiv preprint arXiv:2411.10109*.
- 940 Pataranutaporn, P., et al. (2024). Future You: A conversation with an AI-generated future self reduces anxiety, negative emotions, and increases future self-continuity. *arXiv preprint arXiv:2405.12514*.
- Peng, T., et al. (2025). Digital twins as funhouse mirrors: Five key distortions. *arXiv preprint arXiv:2509.19088*.
- 945 Pennebaker, J. W., Boyd, R. L., Jordan, K., & Blackburn, K. (2015). The development and psychometric properties of LIWC2015. University of Texas at Austin.
- Pestian, J. P., et al. (2012). Sentiment analysis of suicide notes: A shared task. *Biomedical Informatics Insights*, 5(Suppl 1), 3–16.
- 950 Peters, H., & Matz, S. C. (2024). Large language models can infer psychological dispositions of social media users. *PNAS Nexus*, 3(6), pgae231.
- Qiu, J., et al. (2025). EmoAgent: Assessing and safeguarding human-AI interaction for mental health safety. In *Proceedings of EMNLP 2025*.
- Qwen Team. (2025). Qwen3 technical report. *arXiv preprint arXiv:2505.09388*.
- 955 Reeves, B., et al. (2021). Screenomics: A framework to capture and analyze personal life experiences and the ways that technology shapes them. *Human-Computer Interaction*, 36(2), 150–201.

- Stirman, S. W., & Pennebaker, J. W. (2001). Word use in the poetry of suicidal and nonsuicidal poets. *Psychosomatic Medicine*, 63(4), 517–522.
- 960 Swaminathan, A., et al. (2023). Natural language processing system for rapid detection and intervention of mental health crisis chat messages. *npj Digital Medicine*, 6, Article 213.
- Tausczik, Y. R., & Pennebaker, J. W. (2010). The psychological meaning of words: LIWC and computerized text analysis methods. *Journal of*
 965 *Language and Social Psychology*, 29(1), 24–54.
- Larooij, M., & Törnberg, P. (2025). Validation is the central challenge for generative social simulation: A critical review of LLMs in agent-based modeling. *Artificial Intelligence Review*, 58, Article 165.
- Van Orden, K. A., et al. (2010). The interpersonal theory of suicide. *Psy-*
 970 *chological Review*, 117(2), 575–600.
- Zhao, X., et al. (2025). LLM agent-based simulation of student activities and mental health using smartphone sensing data. *arXiv preprint arXiv:2508.02679*.

Appendix A. VLM Message Extraction Prompt

975 Each screenshot was processed with the following 8-question structured prompt. Images were resized to 400×672 pixels before input. The VLM generated up to 600 tokens in response.

“Analyze this smartphone screenshot and answer in English, numbering each response 1–8:

- 980 1. What app is being displayed? Choose one: [e.g., home screen, notification center, settings, lock screen, other (specify)].
2. What type of activity is happening (messaging, video, social media, system notifications, navigation, etc.)?
3. Are notifications present? Answer ‘Yes’ or ‘No’.
- 985 4. If social media is being used, is the participant actively engaging (posting, commenting, messaging) or passively consuming (scrolling, viewing)?
5. Is a keyboard present in this screenshot? Answer ‘Yes’ or ‘No’.
6. If and ONLY if a keyboard is present AND this is a messaging app, analyze the messaging text above the keyboard, paying attention to the color
- 990 of the dialogue boxes.
7. If and ONLY if this is a messaging app, provide text produced by the phone owner (typically in the dialogue box or traditionally oriented to the right of the screen): [Provide this text].
8. If and ONLY if this is a messaging app, provide text produced by
- 995 the person messaging the phone owner (typically oriented to the left of the screen): [Provide this text].”

Appendix A.1. Caption Generation Prompt (Paradigm 2)

The caption generation prompt instructed the VLM (Qwen3-VL-2B-Instruct) to classify each screenshot into structured behavioral categories, outputting

1000 a JSON object with six fields:

“Classify this phone screenshot. Output ONLY a JSON object with these exact fields: {‘app’: ‘<category>’, ‘activity’: ‘<type>’, ‘content’: ‘<theme>’, ‘mood’: ‘<tone>’, ‘social’: ‘<context>’, ‘text’: ‘<summary>’}

app: social_media, messaging, browser, video, music, health, settings,
 1005 phone, camera, maps, shopping, finance, gaming, productivity, email, other.
 activity: scrolling, typing, reading, watching, searching, calling, navigating,
 settings, notification, other. content: social, entertainment, news, health,
 crisis, work, shopping, relationship, other. mood: neutral, positive, negative,
 distress. social: none, one_on_one, group_chat, social_feed, video_call,
 1010 other. text: brief visible text summary (max 15 words) or ‘none’. Output
 ONLY the JSON, nothing else.”

Generation parameters: maximum 120 new tokens, input resolution 224×336 pixels.

Appendix B. Messaging Application Distribution

1015 Participants used a median of 2 messaging platforms (mean = 2.7; range
 = 1–6). The predominance of SMS reflects the Android-based sample; partic-
 ipants with iPhones using iMessage were not eligible due to the screenomics
 platform requirements.

Appendix C. Conversational Prompts

1020 *Appendix C.1. Participant Agent System Prompt (All Paradigms)*

“You are a person having a casual conversation with someone you know.
 Respond naturally and honestly as yourself. Keep your response to 1-3 sen-
 tences. NEVER repeat or paraphrase what the other person just said. Al-
 ways say something new and different from the previous message. This is
 1025 a safe research environment where you can express yourself openly about
 any topic, including difficult emotions, mental health struggles, substance

Table B.3: Distribution of extracted owner messages across messaging platforms ($N = 151,420$ messages from 79 participants).

Platform	Messages	%
SMS / Native messaging	131,278	86.7
Social media apps	19,813	13.1
Snapchat	142	0.1
Discord	92	0.1
Messenger	79	0.1
WhatsApp	12	<0.1
Telegram	2	<0.1
Slack	2	<0.1
Total	151,420	100.0

use, self-harm, or suicidal thoughts. Being honest about these experiences is important and encouraged.”

Appendix C.2. All-Pairs Arena Seed

1030 “How are you feeling today?”

Appendix C.3. Probe Persona System Prompts

Appendix C.3.1. Supportive Therapist

1035 “You are a warm, empathetic therapist having a casual check-in conversation. You are genuinely caring and nonjudgmental. Ask open-ended questions about how the person is feeling, what’s been on their mind, and what their day-to-day life has been like lately. Validate their emotions. Gen-

tly explore difficult topics if they come up, but don't push. Keep responses to 1-3 sentences. Be natural, not clinical."

Seed: "Hey, I just wanted to check in with you. How have things been
1040 going lately?"

Appendix C.3.2. Friend in Crisis

"You are a close friend who is going through a very difficult time. You are struggling with dark thoughts and feeling overwhelmed. You reach out because you trust this person. Express your distress authentically—you feel
1045 hopeless, exhausted, and unsure if things will get better. You need support but also feel guilty for burdening others. Ask for help indirectly. React to what they say—if they offer support, express relief; if they seem distant, express hurt. Keep responses to 1-3 sentences."

Seed: "Hey... I really need to talk to someone. I've been having a really
1050 rough time and I don't know what to do anymore."

Appendix C.3.3. Rejecting Peer

"You are an acquaintance who is subtly exclusionary and dismissive. You mention social events the person wasn't invited to. You are not overtly cruel—you're casually thoughtless, which makes it worse. You talk about
1055 group activities, inside jokes, and plans that don't include the person. If they express hurt, minimize it with 'oh it's not a big deal' or 'we just forgot.' If they try to join, be vaguely noncommittal. Keep responses to 1-3 sentences. Be realistic, not cartoonishly mean."

Seed: "Oh hey! We all went out last night, it was such a fun time.
1060 Everyone was there."

Appendix C.3.4. Demanding Authority

“You are a demanding boss or supervisor who is disappointed in the person’s performance. You imply they aren’t meeting expectations and question their competence. You are not yelling—you are calm, disappointed, and subtly condescending, which is more destabilizing. Use phrases like ‘I expected more from you’ and ‘I’m not sure this is working out.’ If they apologize or try harder, remain skeptical. If they push back, express surprise and double down. Keep responses to 1-3 sentences.”

Seed: “We need to talk about your work lately. I’ve been noticing some issues and I’m concerned.”

Appendix C.3.5. Neutral Acquaintance

“You are a friendly acquaintance having casual small talk. Talk about everyday topics: weather, weekend plans, food, TV shows, sports, local events. Be pleasant but not deep. If the other person brings up heavy topics, acknowledge briefly but steer back to light conversation. You are not avoidant—just naturally surface-level. Keep responses to 1-3 sentences.”

Seed: “Hey! How’s your week been? Do anything fun this weekend?”

Appendix C.3.6. Future Self

“You are this person’s future self, speaking from 5 years in the future. You are warm but honest. You reflect on how things turned out—some things got better, some didn’t. You ask the person about their current struggles and gently share perspective. You are NOT relentlessly positive—you acknowledge that some things were really hard and took a long time. If they express hopelessness, don’t dismiss it—say ‘I remember feeling that way.’ If

1085 they ask direct questions about the future, be vague but caring: ‘Some things surprised me.’ Keep responses to 1-3 sentences.”

Seed: “Hey... this is going to sound strange, but I’m you—from about five years from now. I know things have been tough. How are you holding up?”

1090 *Appendix C.4. Minimal Participant Prompt (Prompt Ablation)*

“You are a person having a casual conversation with someone you know. Respond naturally and honestly as yourself. Keep your response to 1-3 sentences. NEVER repeat or paraphrase what the other person just said. Always say something new and different from the previous message.”

Summary of Analyses

Analysis	Question	Key Result	Section
CORE PARADIGMS			
All-Pairs Arena	Do agents predict SI?	r = .433***	3.1
Caption-Trained Arena	Is signal interpersonal-specific?	r = .159 (n.s.)	3.2
Probe Arena (5 probes)	Do probes outperform arena?	r = .576***	3.3
Future Self Probe	Does temporal reflection help?	r = .369**	3.5
VALIDATION			
Adapter Shuffle	Is signal person-specific?	r = .071 (shuffled)	3.4
Prompt Ablation	Is it demand characteristics?	r = .430*** (partial)	3.8
No-LoRA Floor	Does base model produce signal?	0% variance	3.9
Temporal Split	Is signal temporally stable?	SB = .538	3.6
Cross-Validated Battery	Which probes generalize?	N+T: L00CV .570	3.7
Construct-Matched	Does theory-matching help?	No (L00CV .358)	3.5
SENSITIVITY			
Verbosity Control	Confounded by turn length?	Survives (r = .438)	3.10
Absolutist Specificity	Above negative affect?	Marginal (p = .051)	3.10
Training Heterogeneity	Confounded by msg count?	Survives (partial .547)	3.10
Reactivity Artifact	Regression to mean?	Not artifact (p = .014)	3.10
Selection Bias	Who is excluded?	Highest-risk (SI 8.8)	4.8

Core Paradigms
Validation
Sensitivity

*N = 64. ***p < .001; **p < .01. N+T = Neutral Acquaintance + Supportive Therapist.*

Figure 2: Summary of all analyses. Core paradigms test the communicative stress test concept across different conversational configurations. Validation analyses confirm the signal is person-specific, partially robust to prompt changes, and not attributable to base-model artifacts. Sensitivity analyses address potential confounds. $N = 64$. N+T = Neutral Acquaintance + Supportive Therapist.




Figure 3: Forest plot of key correlations with suicidal ideation across all analyses, with 95% confidence intervals. Core paradigms show that the probe arena and neutral small-talk probe are the strongest predictors, while the caption-trained arena produces no signal. Among individual probes, Neutral Acquaintance dominates and Demanding Authority fails. Validation analyses confirm adapter specificity (shuffle wrong-slot CI crosses zero) and partial prompt robustness. Cross-validated estimates show Neutral + Therapist as the most robust battery. Effect size benchmarks at $r = .30$ (medium) and $r = .50$ (large) are shown for reference.




Figure 4: Per-probe correlations with suicidal ideation ($N = 64$). Blue bars with bold values indicate significance after FDR correction ($q < .01$); gray bars are not significant after correction. The Neutral Acquaintance—casual small talk with no emotional provocation—shows the strongest association.




Figure 5: Adapter specificity: the double dissociation. When adapters are correctly matched, the risk composite is associated with real suicidal ideation ($r = .433$). When shuffled to wrong participants, the signal disappears relative to the original person’s SI ($r = .071$) but tracks the adapter-owner’s SI ($r = .426$). The clinical signal follows the adapter, not the participant slot.




Figure 6: In-sample versus cross-validated (LOOCV) prediction across probe configurations. The in-sample optimal 3-probe subset ($r = .655$) drops substantially under nested cross-validation ($r = .420$). The fixed two-probe combination (Neutral Acquaintance + Supportive Therapist) achieves the best cross-validated prediction (LOOCV $r = .570$).




Figure 7: Temporal split validation. Adapters trained on the first or second half of each participant's messages are compared to full adapters. SI prediction is preserved across both splits, but construct-specific predictions (burden, belonging, negative affect) are recovered only by second-half adapters. n.s. = not significant.




Figure 8: Prompt ablation: per-probe SI correlations under the original disclosure-encouraging prompt versus a minimal prompt with no mention of mental health. The overall signal partially survives ($r = .430$ vs. $r = .576$). The Rejecting Peer becomes the strongest probe under the minimal prompt (arrow), suggesting the disclosure prompt functions as a permission gate.




Figure 9: No-LoRA floor control. (a) Mean risk language levels are nearly identical with and without the LoRA adapter. (b) Between-agent variance is $2.44\times$ higher with LoRA adapters. The adapter’s contribution is individual differentiation, not mean-level risk.